

# HOMEWORK 1: SEQUENTIAL ALGORITHMS

*DEADLINE: OCTOBER 24*

1. Implement in the C programming language the following algorithms:
  - a. Matrix multiplication;
  - b. Gaussian elimination algorithm for solving a linear system of equations assuming that the pivot element cannot be zero: [http://en.wikipedia.org/wiki/Gaussian\\_elimination](http://en.wikipedia.org/wiki/Gaussian_elimination);
  - c. Dijkstra's algorithm for computing the shortest path between nodes in a graph stored as an matrix: [https://en.wikipedia.org/wiki/Dijkstra%27s\\_algorithm](https://en.wikipedia.org/wiki/Dijkstra%27s_algorithm);
  - d. Sieve of Eratosthenes for finding all prime numbers to a given limit: [https://en.wikipedia.org/wiki/Sieve\\_of\\_Eratosthenes](https://en.wikipedia.org/wiki/Sieve_of_Eratosthenes);
2. Implement in the C programming language the following sorting algorithms:
  - a. Bubble sort: [http://en.wikipedia.org/wiki/Bubble\\_sort](http://en.wikipedia.org/wiki/Bubble_sort);
  - b. Bucket sort: [http://en.wikipedia.org/wiki/Bucket\\_sort](http://en.wikipedia.org/wiki/Bucket_sort);
  - c. Counting sort: [http://en.wikipedia.org/wiki/Counting\\_sort](http://en.wikipedia.org/wiki/Counting_sort);
  - d. Insertion sort: [http://en.wikipedia.org/wiki/Insertion\\_sort](http://en.wikipedia.org/wiki/Insertion_sort);
  - e. Selection sort: [http://en.wikipedia.org/wiki/Selection\\_sort](http://en.wikipedia.org/wiki/Selection_sort);
  - f. Quick sort: <https://en.wikipedia.org/wiki/Quicksort>.
3. Initialise the algorithms with uniformly distributed random numbers;
4. Choose for each algorithm one large problem size (i.e. array dimension) and execute it on the `stud1.itec.aau.at` parallel machine using the Slurm workload manager;
5. Use the GNU gprof profiler to measure the execution time of each algorithm;
6. Report and explain the results in a simple PDF file.

Important requirements:

1. Measure only the core execution time of each algorithm without random number generation, array/matrix initialisation, and any I/O operations.
2. Declare the main array or matrix data structures as static global variables. Do not use dynamic memory allocation using `malloc`.
3. Implement the algorithms as short (few lines) and as simple as possible, focused on the core functionality (e.g. without safety checks, small optimisations, or redundant tests).